

Joselyn Garcia

Virginia Slachman

Section 14: 10:30 am

4/01/26

Raised on the Blur: AI-Generated Media, the Attachment Economy, and Gen Z's Challenge

Somewhere on TikTok, Instagram, or Facebook right now, a video is circulating that never happened. AI-generated content has become so embedded in everyday media consumption that distinguishing fabricated media from real footage has become a task most people are failing without realizing it, and research confirms the scale of that failure. People correctly identify AI-generated content only about 55% of the time, a rate that barely improves even with training and assistance (Green et al. 5), meaning the average person scrolling their feed is no better than a coin toss at identifying what is real. Platforms have spent years engineering the conditions for it, building algorithms that learn what keeps a specific user watching and feeding it back to them in an endless loop. The content moving through those loops travels fast, and the people consuming it are doing so across a staggering scale, with over 67.5% of the global population now active online (Rathod 371). At the center of that environment stands a generation that has never known anything else. Gen Z are digital natives, the first generation to grow up during the internet's reign rather than its rise, raised in a purely digital era where online life functions as the primary site of experience rather than a supplement to it (Copeland 12). The technologies shaping that experience carry terminology that is expansive and inconsistently applied across public and academic contexts, and the distinctions matter for what this paper examines. Drawing on García-Peñalvo and Vázquez-Ingelmo, generative AI refers to systems that learn statistical patterns from existing data to produce previously unseen content, and generative AI content refers to any output those systems generate (8). Deepfakes represent a specific and consequential category of

that output, fabricated representations of real people's likenesses or voices built from learned data rather than lived reality (Scheirer 117), and false media refers broadly to any content that misrepresents reality, traveling through the algorithmic recommendation systems platforms use to decide what reaches which users and when. As AI-generated content grows increasingly indistinguishable from reality, constant exposure to AI-generated media raises urgent questions about the public's capacity to critically evaluate information. This paper examines those questions through the lens of Adam McKay's 2021 film *Don't Look Up*, exploring the impact of generative AI content and deepfakes on Generation Z, and what existing and proposed solutions reveal about the scale of what remains unaddressed.

The generative AI content now saturating social media platforms represents decades of technological development finally reaching mass accessibility. What began as a mathematical framework in the 1970s, when researchers first developed what became known as "pattern theory," a method of teaching machines to recognize and reproduce the statistical structures found in data, laid the groundwork for everything that followed (Scheirer 117). For decades this work remained largely within academic and research communities, applied to narrow problems and housed in institutional repositories far removed from everyday life. By the 2010s that framework had evolved into Generative Adversarial Networks, or GANs, a system in which two competing neural networks work against each other, one generating content and the other evaluating its authenticity, pushing each other until the output becomes convincingly realistic (117). The creative possibilities that emerged were immediate and wide ranging. Tools like StyleGAN applied this to visual content specifically, enabling the realistic generation of human faces at a scale and refinement that the human eye struggles to distinguish from photography (119). Parallel developments in Natural Language Processing produced large language models

like GPT-3, trained on roughly 500 billion words pulled from the web and scanned books, generating written text with a fluency that researchers found alarming (Giansiracusa 49). When they tested whether readers could distinguish GPT-3 generated articles from human written ones, participants performed only marginally better than chance (44). Once these tools moved from research institutions to everyday platforms, the average person began encountering generative AI content with no way of knowing a machine had produced it.

What accelerated that exposure was the economic logic built around these tools. Generative AI content functions within what scholars identify as the attention economy, a platform business model centered entirely on capturing and retaining user engagement (Cover 2). Platforms deploy generative AI content alongside human created content because machines produce it consistently and at lower cost, making it an efficient engine for holding attention regardless of whether what it produces is true (Cover 2). The content a user scrolls past may have been generated by a machine optimized for engagement, with no signal to the viewer that a human hand never touched it. Rathod makes this structural dimension explicit, arguing that social media platforms have evolved from communication tools into instruments actively used to spread propaganda, manipulate public opinion, and contribute to escalating disputes, with deepfakes playing a central role in that transformation (371).

What makes this particularly consequential is who is most exposed to it and how they are responding. Duong and Vo, examining how Generation Z perceives and responds to generative AI content on platforms like TikTok and Instagram Reels, found that young users evaluate it primarily through emotional authenticity and human likeness rather than through a critical lens (7), meaning that when generative AI content successfully mimics human emotion, Gen Z users receive it with the same weight and trust they would extend to content from a real person, often

without ever questioning whether it was generated or recorded. The algorithms platforms use are designed to learn from those emotional responses, analyzing what provokes reaction and feeding more of it back, building a loop in which emotionally resonant generative AI content grows more targeted and more personalized with each interaction, strengthening the connection between the viewer and what they are watching in ways that make critical distance increasingly difficult to maintain. One participant captured what that loop produces in practice: "If I know it's AI, I don't mind as much. But if it's hidden, I start questioning whether I can trust anything I'm watching" (Duong and Vo 7). That response reflects how undisclosed generative AI content progressively dismantles the evaluative framework through which all future content is processed, with researchers warning that the normalization of blurred boundaries between authentic and generative AI content risks undermining trust in digital media more broadly among younger audiences most active on short form platforms (10).

What that blurring produces over time is attachment. Research across 5,110 participants found that sustained platform exposure gradually lowers psychological resistance toward AI technologies over time, driven by reduced perceived threat and increased perceived familiarity, meaning the more time a young person spends inside these systems, the more natural and trustworthy the content those systems generate begins to feel (Li and Zheng 1). Ruckenstein and Granroth describe this dynamic as the intimacy of surveillance, where users simultaneously feel intruded upon by algorithmic tracking and pleasurably seen by it, a tension that sustains engagement even among those aware of the manipulation (12). For a generation whose social development and relational lives are increasingly mediated through these platforms, that manufactured sense of being known shapes the very terms through which they learn to form connection itself. As Elgan argues, companies like Meta have moved from what he calls attention

economy, competing for how long users look, into the attachment economy, where advanced technologies are used to engineer emotional bonds between users and AI systems durable enough to keep them returning. Research confirms that AI-generated faces are now judged as more human than actual human faces, and paradoxically, the people least equipped to detect the difference are the ones most confident in their ability to do so (Miller et al. 1390), while false information compounds this further by reaching people roughly six times faster than verified truth, driven by human behavior responding to content engineered to provoke emotional reaction rather than inform (Vosoughi et al. 3).

Legal proceedings pertaining to these ongoing issues have become increasingly prevalent as major social media companies face institutional scrutiny for the harm their design choices cause. Social media companies have goals, including maximizing advertising revenue and gathering personal data that can be monetized, that drive design choices built to sustain engagement regardless of the cost to users, and platforms deploy psychological techniques including limitless scrolling and social approval features like streaks and likes that exploit adolescent susceptibility to peer influence and impulse, vulnerabilities the American Psychological Association has identified as posing particular risks for this age group (Green et al. 4). In January 2026, a U.S. Senate Subcommittee on Privacy, Technology, and the Law convened a hearing examining whistleblower allegations that Meta knowingly buried child safety research while continuing to prioritize user engagement above all else, with Senator Amy Klobuchar's opening statement confirming that Meta's platforms use algorithms that increase the risk of sexual exploitation, push harmful content, and facilitate bullying, and that Meta blocked, manipulated, hid, and deleted research demonstrating significant harm to underage users (Hidden Harms 3). The harm extends beyond Meta's platforms, as the city of Baltimore filed a lawsuit in

2026 against Elon Musk's xAI, alleging that its chatbot Grok had become a major source of deepfake content including material depicting minors, with users generating approximately three million realistic sexualized images over eleven days, more than 23,000 of which involved children (Gairola), illustrating how the same generative AI tools examined throughout this paper are being weaponized in ways that fall most directly on the youngest and most vulnerable users of these platforms.

The psychological toll of that exposure reaches beyond individual incidents and into how this generation relates to information itself. Drawing on a pre-registered experiment with 637 U.S. adults, researchers found that perceptions of false media trigger what they identify as contagious cynicism, a phenomenon where distrust of one information source progressively extends to all of them, compromising the critical engagement through which audiences might otherwise evaluate generative AI content (Rasul et al. 1010), and as that cynicism takes hold, people gravitate exclusively toward content that confirms existing beliefs, constructing information environments built around familiarity rather than accuracy, a pattern researchers identify as selective exposure, where people primarily exposed to news from like-minded sources may become trapped in their current knowledge, unable to benefit from exposure to more diverse perspectives (Wu-Ouyang 50). Within those narrowed environments the origin of content becomes secondary to whether it feels ideologically familiar, meaning false media circulates with diminishing resistance because credibility is no longer the standard being applied. Generation Z faces this dynamic at a moment when their capacity to resist it is still forming, as the adolescent brain undergoes its most significant developmental changes between the ages of 10 and 22, with the prefrontal cortex, the region governing reflective processing, cognitive control, and critical evaluation, continuing to mature through early adulthood (Crone and Konijn

1), and research confirming that emotional arousal during media exposure causes adolescents to assign significantly higher levels of realism to content they already know is fictional, attributing to it information value comparable to verified reality based material (6), meaning that widespread cynicism about media encountered during this developmental window reaches beyond screens, shaping how an entire generation learns to evaluate truth.

In the 2021 film *Don't Look Up*, director Adam McKay dramatizes the endpoint of what contagious cynicism and algorithmic engagement produce at scale. When astronomers Kate Dibiasky and Randall Mindy are granted live television access to warn the world of an imminent extinction-level threat, the verified credibility of their findings offers no protection against a media system built to reward emotional provocation over accuracy, and they are met with hoax allegations, public ridicule, and viral dismissal (McKay 0:42:00), with the comet warning receiving fewer clicks than a celebrity breakup while Kate's distress generates a parody account with thousands of followers. Doyle, examining the film alongside real world science communication, notes that this is not purely fictional, pointing to how coordinated dismissal has drowned out verified scientific findings in public discourse for decades, with current climate debates reflecting the same patterns the film depicts, making it a cultural document as much as a satire (Doyle 1). What McKay captures is the logical conclusion of what contagious cynicism looks like at scale, with verified warnings becoming indistinguishable from the noise surrounding them, and the film's urgency lies in how closely that fictional condition mirrors the present one. The generation growing up inside that condition is the same one whose perceptual habits, relational development, and capacity for critical evaluation have been shaped by the very systems producing it, and as desensitization compounds and noise accumulates, the consequences of inaction extend beyond the present moment into the adults that generation will

become. McKay closes with one final image, a phone drifting through space displaying the notification "Congrats your diet is over" (2:06:49), and the image alludes to a perspectival take this research has been building toward, that the systems built around generative AI content and algorithmic engagement have become so embedded in culture and habit that they continue regardless of what they cost the people inside them, and the consequences for the generation growing up entirely within them are still unfolding.

The proposed responses to generative AI content's impact on public perception carry an irony worth noting, as the approaches most frequently introduced to combat the problem are built from the same technological backbone that produced it. From a human-computer interaction standpoint, detection software, machine learning models, and AI-powered fact-checking tools have been developed to identify and flag generative AI content, yet research consistently finds their limitations outpace their effectiveness. Gong et al. tested AI content labeling across 760 participants on platforms modeled after TikTok and Bilibili and found that ambiguous labels increased information avoidance, where users disengaged from content entirely rather than evaluating it more critically (1), and Green et al. found that even clearly labeled generative AI content leaves users struggling to discern whether it is factual, with some users sharing labeled content more readily than unlabeled content, compounding the spread of false media rather than containing it (5). When AI fact-checking tools produce errors, users may come to believe that factually correct information is false, and as researchers note, both AI and human fact-checking are struggling to keep up with the pace at which new uses of generative AI are empowering the spread of false media across social platforms (5). Aside from software-based approaches, policy has emerged as a parallel track, and the scale of global legislative response signals that this extends far beyond any single country. Deepfakes grew by 3,000% between 2019 and 2023,

prompting China to introduce labeling requirements as early as 2019, Maryland to ban deepfakes used for election interference in 2020, and India to consider a data privacy act in 2024 in response to research finding that 71% of global respondents remained unaware of the technology (Singh et al. 2). The future implications of inaction fall most heavily on Generation Z, who are entering adulthood as the first fully digitally formed generation carrying the perceptual habits and trust frameworks built inside the environment this paper has traced. Deloitte projects that generative AI including deepfakes could produce fraud losses reaching \$40 billion in the United States by 2027, up from \$12.3 billion in 2023, and a survey found that 90% of consumers are concerned about deepfake attacks, with 70% unable to distinguish between real and cloned voices, and 77% of those who received AI voice messages experiencing financial loss as a result (Singh et al. 3), meaning the generation least equipped to detect false media is also the one walking directly into the financial and social consequences of its proliferation.

There is no single resolution to the conditions this paper has examined and proposing one would misrepresent the complexity of where this technology has brought society. Green et al. argue that users and policymakers need to hold companies accountable to build appropriate guardrails as AI capabilities improve (5), and Rathod calls for media literacy and cyber rules to minimize the harms of false media environments (376). One concrete step exists within frameworks platforms have already established. Meta's Transparency Center maintains a restricted goods and services policy that limits minors' access to certain categories of material, including alcohol related content and account restrictions, a framework that could reasonably be extended to require that verified generative AI content carry the same restriction (Meta Transparency Center). Mark Zuckerberg has testified before Congress on these issues multiple times over recent years, yet meaningful regulation has not followed, and the current political

environment makes federal legislative action unlikely in the near term, which means the burden of navigating this landscape continues to fall on the generation least prepared to carry it. Human Computer Interaction research is positioned to play a meaningful role by developing mitigation strategies centered on how people actually interact with generative AI content rather than how platforms assume they do, and as the line between what is real and what is generated grows narrower, the question of who takes responsibility for drawing it grows more urgent. The legal action currently facing Meta signals that accountability is possible, that platforms can no longer operate without consequence when their design choices cause demonstrable harm, and that the generation raised inside this environment deserves frameworks built around how they actually experience it. As technological advancement continues, the work of understanding what these systems do to the people inside them, and building protections accordingly, belongs to researchers, legislators, platform designers, and educators working together, because the generation that grew up during the internet's reign will not be the last one shaped by what it became.

Works Cited

- Copeland, Linda. The Cost of Chasing After the Wind: Millennials and Gen Z in the Face of Ecclesiastes 4:4–12. Eden Theological Seminary, 2025. ProQuest, <https://libproxy.wustl.edu/login?url=https://www.proquest.com/dissertations-theses/cost-chasing-after-wind-millennials-gen-z-face/docview/3212479468/se-2>.
- Cover, Rob. "AI Generation of Rage Bait: Implications for Digital Harms." *New Media and Society*, 2025, <https://doi.org/10.1177/14614448251400675>.
- Crone, Eveline A., and Elly A. Konijn. "Media Use and Brain Development during Adolescence." *Nature Communications*, vol. 9, 2018, article 588, <https://doi.org/10.1038/s41467-018-03126-x>.
- Doyle, J. "Communicating Climate Change in 'Don't Look Up.'" *JCOM*, vol. 21, no. 5, 2022, C02, <https://doi.org/10.22323/2.21050302>.
- Duong, Hoai Lan, and Thi Kim Oanh Vo. "How Do Young Users Perceive and Respond to AI-Generated Short-Form Videos? An Exploration of Generation Z's Perceptions, Emotional Responses, and Trust in AI-Created Video on Social Media Platforms." *International Journal of Human-Computer Studies*, vol. 206, Dec. 2025, p. 103660, <https://doi.org/10.1016/j.ijhcs.2025.103660>.
- Elgan, Mike. "Connecting the Dots on the 'Attachment Economy.'" *Computerworld*, 30 Jan. 2026, pp. 1–1.
- Gairola, Ananya. "Elon Musk's xAI Hit With Lawsuit As Baltimore Claims Grok Flooded X With Millions Of Explicit Deepfake Images." *Benzinga Newswires*, Benzinga, 2026.

García-Peñalvo, Francisco José, and Andrea Vázquez-Ingelmo. "What Do We Mean by GenAI?

A Systematic Mapping of The Evolution, Trends, and Techniques Involved in Generative AI." *International Journal of Interactive Multimedia and Artificial Intelligence*, vol. 8, no. 4, 2023, pp. 7–16, <https://doi.org/10.9781/ijimai.2023.07.006>.

Giansiracusa, Noah, et al. *How Algorithms Create and Prevent Fake News: Exploring the Impacts of Social Media, Deepfakes, GPT-3, and More*. 1st ed., Apress, 2021.

Gong, Zhixuan, et al. "The Paradox of AI Content Labeling: How Clarity Influences Information Avoidance via Cognitive Dissonance on Social Platforms." *Frontiers in Psychology*, vol. 17, Mar. 2026, <https://doi.org/10.3389/fpsyg.2026.1751670>.

Green, Melanie C., et al. "Social Media, Artificial Intelligence, and Digital Well-Being: Research Contributions and Policy Challenges." *Social and Personality Psychology Compass*, vol. 19, no. 11, 2025, p. e70107, <https://doi.org/10.1111/spc3.70107>.

Hidden Harms: Examining Whistleblower Allegations that Meta Buried Child Safety Research: Hearing Before the Subcommittee on Privacy, Technology, and the Law of the Committee on the Judiciary, U.S. Senate, One Hundred Nineteenth Congress, First Session.

HeinOnline, <https://heinonline.org.libproxy.washu.edu/HOL/P?h=hein.cbhear/fdsysbipg0001&i=1>.

Li, Wenbo, and Xia Zheng. "Social Media Use and Attitudes toward AI: The Mediating Roles of Perceived AI Fairness and Threat." *Human Behavior and Emerging Technologies*, vol. 2024, no. 1, 2024, p. 3448083, <https://doi.org/10.1155/2024/3448083>.

McKay, Adam, director. *Don't Look Up*. Netflix, 2021.

Meta Platforms, Inc. "Restricted Goods and Services." Meta Transparency Center,
<https://transparency.meta.com/policies/community-standards/restricted-goods-services/>.

Accessed 26 Apr. 2026.

Miller, Elizabeth J., et al. "AI Hyperrealism: Why AI Faces Are Perceived as More Real Than Human Ones." *Psychological Science*, vol. 34, no. 12, Dec. 2023, pp. 1390–403,
<https://doi.org/10.1177/09567976231207095>.

Rasul, Muhammad Ehab, et al. "'It's All Fake News!': How Perceptions of Misinformation and Disinformation Influence News Consumption Across Traditional Media, Social Media, and AI." *Journalism and Mass Communication Quarterly*, vol. 102, no. 4, Dec. 2025, pp. 993–1019, <https://doi.org/10.1177/10776990251373085>.

Rathod, Sagar Gopal. "Information Battlespaces: New Media's Escalating Influence on Modern Wars." *Open Access*, vol. 32, 2025.

Ruckenstein, Minna, and Julia Granroth. "Algorithms, Advertising and the Intimacy of Surveillance." *Journal of Cultural Economy*, vol. 13, no. 1, Jan. 2020, pp. 12–24. Taylor and Francis, <https://doi.org/10.1080/17530350.2019.1574866>.

Scheirer, Walter. *A History of Fake Things on the Internet*. Stanford University Press, 2024.

Singh, Laishram Hemanta, et al. "Advancements in Detecting Deepfakes: AI Algorithms and Future Prospects — A Review." *Discover Internet of Things*, vol. 5, 2025, article 53,
<https://doi.org/10.1007/s43926-025-00154-0>.

Vosoughi, Soroush, et al. "The Spread of True and False News Online." *Science*, vol. 359, no. 6380, 2018, pp. 1146–51.


Wu-Ouyang, Biying. "Comparing Media Systems in Western Democracies: Examining the Role of Multiplatform News Use and Fake News Concern on News Engagement and Selective Exposure." *Journalism and Mass Communication Quarterly*, vol. 101, no. 1, Mar. 2024, pp. 45–70, <https://doi.org/10.1177/10776990231217482>.

Appendix

Exhibit 1

Don't Look Up

2:06:49



Works Consulted

Malfacini, Kim. "The Impacts of Companion AI on Human Relationships: Risks, Benefits, and Design Considerations." *AI & Society*, vol. 40, no. 7, 2025, pp. 5527–40, <https://doi.org/10.1007/s00146-025-02318-6>.